

# Apply different Architectures on MNIST dataset using Keras

In [19]:

```
import warnings
warnings.filterwarnings("ignore")
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use th
is command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
#from keras.utils.visualize_util import to_graph
from keras.models import Sequential
#to_graph(Sequential())
```

## READING DATA

In [4]:

```
%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [5]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
 11493376/11490434 [=====] - 164s 14us/step

In [6]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%
d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d
, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)  
 Number of training examples : 10000 and each image is of shape (28,28)

In [7]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [8]:

```
# after converting the input images from 3d to 2d vectors
print("Number of training examples :", X_train.shape[0], "and each image is of shape(%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape(784)

Number of training examples : 10000 and each image is of shape (784)

In [9]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 
X_train = X_train/255
X_test = X_test/255
```

In [10]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])
# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs
Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

In [11]:

```
# some model parameters
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 112
nb_epoch = 20
print(input_dim)
```

784

## Model 1 -- with 2 Hidden layers

# 1. MLP + ReLU + adam

In [22]:

```
from keras.layers import Activation, Dense
```

In [23]:

```
# Multilayer perceptron
# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(\theta, \sigma) = N(\theta, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(\theta, \sigma) = N(\theta, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(610, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(325, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))
model_relu.summary()
```

WARNING:tensorflow:From C:\anaconda\lib\site-packages\tensorflow\python\framework\op\_def\_library.py:263: colocate\_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

Layer (type)	Output Shape	Param #
=====		
dense_1 (Dense)	(None, 610)	478850
dense_2 (Dense)	(None, 325)	198575
dense_3 (Dense)	(None, 10)	3260
=====		
Total params: 680,685		
Trainable params: 680,685		
Non-trainable params: 0		

In [24]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

WARNING:tensorflow:From C:\anaconda\lib\site-packages\tensorflow\python\ops\math\_ops.py:3066: to\_int32 (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 19s 322us/step - loss: 0.2113 - acc: 0.9366 - val\_loss: 0.1056 - val\_acc: 0.9668

Epoch 2/20

60000/60000 [=====] - 17s 275us/step - loss: 0.0781 - acc: 0.9759 - val\_loss: 0.0785 - val\_acc: 0.9743

Epoch 3/20

60000/60000 [=====] - 17s 276us/step - loss: 0.0468 - acc: 0.9848 - val\_loss: 0.0803 - val\_acc: 0.9747

Epoch 4/20

60000/60000 [=====] - 16s 266us/step - loss: 0.0330 - acc: 0.9896 - val\_loss: 0.0759 - val\_acc: 0.9763

Epoch 5/20

60000/60000 [=====] - 16s 265us/step - loss: 0.0261 - acc: 0.9914 - val\_loss: 0.0733 - val\_acc: 0.9788

Epoch 6/20

60000/60000 [=====] - 16s 274us/step - loss: 0.0209 - acc: 0.9931 - val\_loss: 0.0723 - val\_acc: 0.9796

Epoch 7/20

60000/60000 [=====] - 17s 275us/step - loss: 0.0168 - acc: 0.9947 - val\_loss: 0.0880 - val\_acc: 0.9760

Epoch 8/20

60000/60000 [=====] - 13s 215us/step - loss: 0.0162 - acc: 0.9950 - val\_loss: 0.0927 - val\_acc: 0.9786

Epoch 9/20

60000/60000 [=====] - 13s 218us/step - loss: 0.0176 - acc: 0.9943 - val\_loss: 0.0786 - val\_acc: 0.9801

Epoch 10/20

60000/60000 [=====] - 13s 214us/step - loss: 0.0155 - acc: 0.9950 - val\_loss: 0.0808 - val\_acc: 0.9819

Epoch 11/20

60000/60000 [=====] - 13s 212us/step - loss: 0.0094 - acc: 0.9969 - val\_loss: 0.0918 - val\_acc: 0.9777

Epoch 12/20

60000/60000 [=====] - 13s 215us/step - loss: 0.0128 - acc: 0.9955 - val\_loss: 0.0905 - val\_acc: 0.9800

Epoch 13/20

60000/60000 [=====] - 13s 213us/step - loss: 0.0110 - acc: 0.9965 - val\_loss: 0.0876 - val\_acc: 0.9811

Epoch 14/20

60000/60000 [=====] - 13s 210us/step - loss: 0.0089 - acc: 0.9971 - val\_loss: 0.1207 - val\_acc: 0.9748

Epoch 15/20

60000/60000 [=====] - 13s 212us/step - loss: 0.0186 - acc: 0.9940 - val\_loss: 0.1070 - val\_acc: 0.9767

Epoch 16/20

60000/60000 [=====] - 13s 212us/step - loss: 0.0068 - acc: 0.9979 - val\_loss: 0.1123 - val\_acc: 0.9796

Epoch 17/20

60000/60000 [=====] - 13s 215us/step - loss: 0.0093 - acc: 0.9971 - val\_loss: 0.0883 - val\_acc: 0.9813

Epoch 18/20

60000/60000 [=====] - 13s 213us/step - loss: 0.0082 - acc: 0.9974 - val\_loss: 0.1028 - val\_acc: 0.9804

Epoch 19/20

```
60000/60000 [=====] - 13s 209us/step - loss: 0.00  
90 - acc: 0.9973 - val_loss: 0.1004 - val_acc: 0.9804  
Epoch 20/20  
60000/60000 [=====] - 13s 209us/step - loss: 0.00  
92 - acc: 0.9971 - val_loss: 0.1196 - val_acc: 0.9782
```

In [25]:


```
#Evaluate your model with accuracy and plot of (NUmber of epoches VS train_and_val_Los
s)
#Train accuracy
score = model_relu.evaluate(X_train, Y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n*****\n')
#test accuracy
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbse=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Train score: 0.01175171572082836  
Train accuracy: 99.61166666666666

\*\*\*\*\*

Test score: 0.11959466045504177  
Test accuracy: 97.82





In [26]:

```
# Weights after training
# 1 2 3
# input->h1->h2->output
w_after = model_relu.get_weights()
# if 2 hidden layer then
# w_after[0]is the input layer weights w_after[1]is the input layer bias weights input
to hidden1
# w_after[2]is the hidde layer weights w_after[3]is the hidde layer bias weights hidden
1to hidden2
# w_after[4]is the hidde layer weights w_after[5]is the hidde layer bias weights hidden
2 to output
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")


ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")

ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")

ax = sns.violinplot(y=out_w, color='g')
plt.xlabel('output layer ')
```

Out[26]:

Text(0.5, 0, 'output layer ')



## 2. MLP + ReLU + adam + batch\_\_normalization

In [27]:

```
from keras.layers.normalization import BatchNormalization
model_batch = Sequential()
model_batch.add(Dense(610, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())
model_batch.add(Dense(325, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())
model_batch.add(Dense(output_dim, activation='softmax'))
model_batch.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_4 (Dense)	(None, 610)	478850
batch_normalization_1 (Batch Normalization)	(None, 610)	2440
dense_5 (Dense)	(None, 325)	198575
batch_normalization_2 (Batch Normalization)	(None, 325)	1300
dense_6 (Dense)	(None, 10)	3260
=====		
Total params: 684,425		
Trainable params: 682,555		
Non-trainable params: 1,870		
=====		

In [28]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 19s 314us/step - loss: 0.18  
04 - acc: 0.9459 - val\_loss: 0.1109 - val\_acc: 0.9638

Epoch 2/20

60000/60000 [=====] - 15s 256us/step - loss: 0.06  
91 - acc: 0.9785 - val\_loss: 0.0789 - val\_acc: 0.9747

Epoch 3/20

60000/60000 [=====] - 15s 258us/step - loss: 0.04  
32 - acc: 0.9864 - val\_loss: 0.0747 - val\_acc: 0.9762

Epoch 4/20

60000/60000 [=====] - 17s 285us/step - loss: 0.03  
20 - acc: 0.9902 - val\_loss: 0.0815 - val\_acc: 0.9748

Epoch 5/20

60000/60000 [=====] - 15s 249us/step - loss: 0.02  
53 - acc: 0.9920 - val\_loss: 0.0845 - val\_acc: 0.9747

Epoch 6/20

60000/60000 [=====] - 15s 249us/step - loss: 0.02  
10 - acc: 0.9934 - val\_loss: 0.0792 - val\_acc: 0.9781

Epoch 7/20

60000/60000 [=====] - 15s 254us/step - loss: 0.02  
08 - acc: 0.9934 - val\_loss: 0.0772 - val\_acc: 0.9770

Epoch 8/20

60000/60000 [=====] - 15s 251us/step - loss: 0.01  
89 - acc: 0.9940 - val\_loss: 0.0783 - val\_acc: 0.9781

Epoch 9/20

60000/60000 [=====] - 15s 247us/step - loss: 0.01  
38 - acc: 0.9950 - val\_loss: 0.0883 - val\_acc: 0.9753

Epoch 10/20

60000/60000 [=====] - 15s 248us/step - loss: 0.01  
47 - acc: 0.9951 - val\_loss: 0.0810 - val\_acc: 0.9782

Epoch 11/20

60000/60000 [=====] - 15s 250us/step - loss: 0.01  
31 - acc: 0.9956 - val\_loss: 0.0811 - val\_acc: 0.9794

Epoch 12/20

60000/60000 [=====] - 17s 275us/step - loss: 0.00  
96 - acc: 0.9968 - val\_loss: 0.0736 - val\_acc: 0.9798

Epoch 13/20

60000/60000 [=====] - 15s 250us/step - loss: 0.00  
96 - acc: 0.9970 - val\_loss: 0.0781 - val\_acc: 0.9798

Epoch 14/20

60000/60000 [=====] - 15s 250us/step - loss: 0.01  
19 - acc: 0.9960 - val\_loss: 0.0736 - val\_acc: 0.9812

Epoch 15/20

60000/60000 [=====] - 15s 255us/step - loss: 0.01  
02 - acc: 0.9965 - val\_loss: 0.0792 - val\_acc: 0.9797

Epoch 16/20

60000/60000 [=====] - 15s 249us/step - loss: 0.00  
81 - acc: 0.9974 - val\_loss: 0.0779 - val\_acc: 0.9791

Epoch 17/20

60000/60000 [=====] - 15s 250us/step - loss: 0.00  
68 - acc: 0.9979 - val\_loss: 0.0846 - val\_acc: 0.9803

Epoch 18/20

60000/60000 [=====] - 15s 246us/step - loss: 0.00  
84 - acc: 0.9972 - val\_loss: 0.0915 - val\_acc: 0.9782

Epoch 19/20

60000/60000 [=====] - 15s 249us/step - loss: 0.00  
72 - acc: 0.9977 - val\_loss: 0.0827 - val\_acc: 0.9816

Epoch 20/20

60000/60000 [=====] - 15s 246us/step - loss: 0.00  
83 - acc: 0.9971 - val\_loss: 0.0799 - val\_acc: 0.9820

In [30]:

```
#Evaluate your model with accuracy and plot of (NUmber of epoches VS train_and_val_los
s)
#Train accuracy
score = model_batch.evaluate(X_train, Y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n***** \n')
#test accuracy
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# List of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```


Train score: 0.0036005233980133805

Train accuracy: 99.88333333333334

\*\*\*\*\*

Test score: 0.07985942602099326

Test accuracy: 98.2



In [31]:

```
# Weights after training
# 1 2 3
# input->h1->h2->output
w_after = model_batch.get_weights()
# if 2 hidden layer then
# w_after[0]is the inpupt layer weights w_after[1]is the input layer bias weights input
to hidden1
# w_after[2]is the hidde layer weights w_after[3]is the hidde layer bias weights hidden
1 to hidden2
# w_after[4]is the hidde layer weights w_after[5]is the hidde layer bias weights hidden
2 to output
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
fig = plt.figure()
plt.title("Weight matrices after model trained")


plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='g')
plt.xlabel('output layer ')
```

Out[31]:

Text(0.5, 0, 'output layer ')



### 3. MLP + ReLU + adam + dropout

In [32]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
from keras.layers import Dropout
model_drop = Sequential()
model_drop.add(Dense(610, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(325, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(output_dim, activation='softmax'))
model_drop.summary()
```

WARNING:tensorflow:From C:\anaconda\lib\site-packages\keras\backend\tensorflow\_backend.py:3445: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.  
Instructions for updating:  
Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 610)	478850
dropout_1 (Dropout)	(None, 610)	0
dense_8 (Dense)	(None, 325)	198575
dropout_2 (Dropout)	(None, 325)	0
dense_9 (Dense)	(None, 10)	3260
Total params: 680,685		
Trainable params: 680,685		
Non-trainable params: 0		

In [33]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=20, verbose=1,
validation_data=(X_test, Y_test))
```



Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 17s 285us/step - loss: 0.73  
05 - acc: 0.8101 - val\_loss: 0.1879 - val\_acc: 0.9447

Epoch 2/20

60000/60000 [=====] - 14s 232us/step - loss: 0.30  
38 - acc: 0.9110 - val\_loss: 0.1403 - val\_acc: 0.9594

Epoch 3/20

60000/60000 [=====] - 14s 231us/step - loss: 0.24  
30 - acc: 0.9283 - val\_loss: 0.1160 - val\_acc: 0.9659

Epoch 4/20

60000/60000 [=====] - 14s 231us/step - loss: 0.20  
77 - acc: 0.9398 - val\_loss: 0.1064 - val\_acc: 0.9686

Epoch 5/20

60000/60000 [=====] - 14s 233us/step - loss: 0.18  
61 - acc: 0.9460 - val\_loss: 0.1031 - val\_acc: 0.9724- lo

Epoch 6/20

60000/60000 [=====] - 14s 233us/step - loss: 0.16  
91 - acc: 0.9505 - val\_loss: 0.0952 - val\_acc: 0.9739

Epoch 7/20

60000/60000 [=====] - 15s 250us/step - loss: 0.15  
59 - acc: 0.9536 - val\_loss: 0.0932 - val\_acc: 0.9730

Epoch 8/20

60000/60000 [=====] - 14s 241us/step - loss: 0.14  
71 - acc: 0.9572 - val\_loss: 0.0946 - val\_acc: 0.9733

Epoch 9/20

60000/60000 [=====] - 14s 234us/step - loss: 0.13  
90 - acc: 0.9604 - val\_loss: 0.0866 - val\_acc: 0.9744

Epoch 10/20

60000/60000 [=====] - 14s 235us/step - loss: 0.13  
31 - acc: 0.9607 - val\_loss: 0.0858 - val\_acc: 0.9766

Epoch 11/20

60000/60000 [=====] - 14s 235us/step - loss: 0.12  
64 - acc: 0.9628 - val\_loss: 0.0799 - val\_acc: 0.9775

Epoch 12/20

60000/60000 [=====] - 14s 235us/step - loss: 0.11  
67 - acc: 0.9654 - val\_loss: 0.0833 - val\_acc: 0.9773

Epoch 13/20

60000/60000 [=====] - 14s 239us/step - loss: 0.11  
27 - acc: 0.9673 - val\_loss: 0.0839 - val\_acc: 0.9748

Epoch 14/20

60000/60000 [=====] - 15s 245us/step - loss: 0.10  
62 - acc: 0.9680 - val\_loss: 0.0743 - val\_acc: 0.9793

Epoch 15/20

60000/60000 [=====] - 14s 237us/step - loss: 0.10  
84 - acc: 0.9691 - val\_loss: 0.0792 - val\_acc: 0.9769

Epoch 16/20

60000/60000 [=====] - 14s 234us/step - loss: 0.10  
08 - acc: 0.9706 - val\_loss: 0.0726 - val\_acc: 0.9796

Epoch 17/20

60000/60000 [=====] - 14s 234us/step - loss: 0.09  
95 - acc: 0.9709 - val\_loss: 0.0782 - val\_acc: 0.9787

Epoch 18/20

60000/60000 [=====] - 15s 248us/step - loss: 0.09  
28 - acc: 0.9732 - val\_loss: 0.0796 - val\_acc: 0.9791

Epoch 19/20

60000/60000 [=====] - 15s 245us/step - loss: 0.09  
34 - acc: 0.9725 - val\_loss: 0.0847 - val\_acc: 0.9771

Epoch 20/20

60000/60000 [=====] - 16s 260us/step - loss: 0.09  
03 - acc: 0.9736 - val\_loss: 0.0811 - val\_acc: 0.9777

In [34]:

```
#Evaluate your model with accuracy and plot of (NUmber of epoches VS train_and_val_los
s)
#Train accuracy
score = model_drop.evaluate(X_train, Y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n***** \n')
#test accuracy
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
fig,ax = plt.subplots(1,1)

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```


Train score: 0.02664580340325483

Train accuracy: 99.25833333333334

\*\*\*\*\*

Test score: 0.08111033427524962

Test accuracy: 97.77




In [35]:

```
w_after = model_drop.get_weights()
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
fig = plt.figure()
plt.title("Weight matrices after model trained")

plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 4. MLP + ReLU + adam + dropout+ batch\_normalization

In [36]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-f
unction-in-keras
from keras.layers import Dropout
model_drop = Sequential()
model_drop.add(Dense(610, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(325, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(output_dim, activation='softmax'))
model_drop.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_10 (Dense)	(None, 610)	478850
batch_normalization_3 (Batch Normalization)	(None, 610)	2440
dropout_3 (Dropout)	(None, 610)	0
dense_11 (Dense)	(None, 325)	198575
batch_normalization_4 (Batch Normalization)	(None, 325)	1300
dropout_4 (Dropout)	(None, 325)	0
dense_12 (Dense)	(None, 10)	3260
=====	=====	=====
Total params: 684,425		
Trainable params: 682,555		
Non-trainable params: 1,870		
=====		

In [37]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=20, verbose=1,  
validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 25s 414us/step - loss: 0.41  
80 - acc: 0.8724 - val\_loss: 0.1510 - val\_acc: 0.9508

Epoch 2/20

60000/60000 [=====] - 20s 334us/step - loss: 0.22  
14 - acc: 0.9328 - val\_loss: 0.1119 - val\_acc: 0.9643

Epoch 3/20

60000/60000 [=====] - 20s 334us/step - loss: 0.17  
71 - acc: 0.9452 - val\_loss: 0.0968 - val\_acc: 0.9698

Epoch 4/20

60000/60000 [=====] - 18s 307us/step - loss: 0.15  
13 - acc: 0.9536 - val\_loss: 0.0906 - val\_acc: 0.9714

Epoch 5/20

60000/60000 [=====] - 19s 319us/step - loss: 0.14  
00 - acc: 0.9577 - val\_loss: 0.0829 - val\_acc: 0.9740

Epoch 6/20

60000/60000 [=====] - 18s 294us/step - loss: 0.12  
27 - acc: 0.9621 - val\_loss: 0.0767 - val\_acc: 0.9758

Epoch 7/20

60000/60000 [=====] - 18s 305us/step - loss: 0.11  
81 - acc: 0.9627 - val\_loss: 0.0730 - val\_acc: 0.9778

Epoch 8/20

60000/60000 [=====] - 19s 310us/step - loss: 0.11  
14 - acc: 0.9651 - val\_loss: 0.0753 - val\_acc: 0.9756

Epoch 9/20

60000/60000 [=====] - 19s 310us/step - loss: 0.10  
29 - acc: 0.9670 - val\_loss: 0.0694 - val\_acc: 0.9774

Epoch 10/20

60000/60000 [=====] - 19s 314us/step - loss: 0.09  
77 - acc: 0.9693 - val\_loss: 0.0640 - val\_acc: 0.9786

Epoch 11/20

60000/60000 [=====] - 18s 306us/step - loss: 0.09  
13 - acc: 0.9720 - val\_loss: 0.0634 - val\_acc: 0.9809

Epoch 12/20

60000/60000 [=====] - 19s 310us/step - loss: 0.08  
56 - acc: 0.9727 - val\_loss: 0.0626 - val\_acc: 0.9806

Epoch 13/20

60000/60000 [=====] - 18s 303us/step - loss: 0.08  
51 - acc: 0.9731 - val\_loss: 0.0653 - val\_acc: 0.9797

Epoch 14/20

60000/60000 [=====] - 18s 303us/step - loss: 0.07  
91 - acc: 0.9746 - val\_loss: 0.0618 - val\_acc: 0.9810

Epoch 15/20

60000/60000 [=====] - 18s 306us/step - loss: 0.07  
44 - acc: 0.9763 - val\_loss: 0.0599 - val\_acc: 0.9819

Epoch 16/20

60000/60000 [=====] - 18s 302us/step - loss: 0.07  
25 - acc: 0.9768 - val\_loss: 0.0594 - val\_acc: 0.9816

Epoch 17/20

60000/60000 [=====] - 18s 304us/step - loss: 0.06  
75 - acc: 0.9783 - val\_loss: 0.0614 - val\_acc: 0.9815

Epoch 18/20

60000/60000 [=====] - 18s 301us/step - loss: 0.06  
77 - acc: 0.9783 - val\_loss: 0.0627 - val\_acc: 0.9804

Epoch 19/20

60000/60000 [=====] - 18s 300us/step - loss: 0.06  
15 - acc: 0.9802 - val\_loss: 0.0601 - val\_acc: 0.9810

Epoch 20/20

60000/60000 [=====] - 19s 310us/step - loss: 0.06  
02 - acc: 0.9807 - val\_loss: 0.0578 - val\_acc: 0.9819

In [38]:

```
#Evaluate your model with accuracy and plot of (NUmber of epoches VS train_and_val_Loss)
#Train accuracy
score = model_drop.evaluate(X_train, Y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n*****\n')
#test accuracy
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```


Train score: 0.014805679358745692

Train accuracy: 99.53666666666666

\*\*\*\*\*

Test score: 0.05782464738052222

Test accuracy: 98.19




In [39]:

```
w_after = model_drop.get_weights()
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
fig = plt.figure()
plt.title("Weight matrices after model trained")

plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## Model 2 -- with 3 Hidden layers

### 1. MLP + ReLU + adam



In [40]:

```
# Multilayer perceptron
# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(\theta, \sigma) = N(\theta, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(\theta, \sigma) = N(\theta, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 
model_relu = Sequential()
model_relu.add(Dense(610, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(420, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(210, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))
model_relu.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_13 (Dense)	(None, 610)	478850
dense_14 (Dense)	(None, 420)	256620
dense_15 (Dense)	(None, 210)	88410
dense_16 (Dense)	(None, 10)	2110
=====	=====	=====
Total params: 825,990		
Trainable params: 825,990		
Non-trainable params: 0		

In [41]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 20s 330us/step - loss: 0.20

76 - acc: 0.9368 - val\_loss: 0.0928 - val\_acc: 0.9695

Epoch 2/20

60000/60000 [=====] - 14s 240us/step - loss: 0.07

94 - acc: 0.9751 - val\_loss: 0.0830 - val\_acc: 0.9755

Epoch 3/20

60000/60000 [=====] - 14s 240us/step - loss: 0.05

22 - acc: 0.9831 - val\_loss: 0.0883 - val\_acc: 0.9752

Epoch 4/20

60000/60000 [=====] - 14s 239us/step - loss: 0.03

83 - acc: 0.9879 - val\_loss: 0.0944 - val\_acc: 0.9699

Epoch 5/20

60000/60000 [=====] - 15s 244us/step - loss: 0.03

26 - acc: 0.9887 - val\_loss: 0.0931 - val\_acc: 0.9754

Epoch 6/20

60000/60000 [=====] - 14s 240us/step - loss: 0.03

06 - acc: 0.9905 - val\_loss: 0.0953 - val\_acc: 0.9740

Epoch 7/20

60000/60000 [=====] - 14s 236us/step - loss: 0.02

56 - acc: 0.9915 - val\_loss: 0.0834 - val\_acc: 0.9781

Epoch 8/20

60000/60000 [=====] - 14s 234us/step - loss: 0.02

44 - acc: 0.9920 - val\_loss: 0.0935 - val\_acc: 0.9783

Epoch 9/20

60000/60000 [=====] - 14s 235us/step - loss: 0.02

32 - acc: 0.9923 - val\_loss: 0.0952 - val\_acc: 0.9763

Epoch 10/20

60000/60000 [=====] - 14s 237us/step - loss: 0.01

34 - acc: 0.9960 - val\_loss: 0.1032 - val\_acc: 0.9778

Epoch 11/20

60000/60000 [=====] - 19s 324us/step - loss: 0.01

75 - acc: 0.9942 - val\_loss: 0.1080 - val\_acc: 0.9745

Epoch 12/20

60000/60000 [=====] - 16s 271us/step - loss: 0.01

68 - acc: 0.9948 - val\_loss: 0.0825 - val\_acc: 0.9803

Epoch 13/20

60000/60000 [=====] - 14s 238us/step - loss: 0.01

86 - acc: 0.9942 - val\_loss: 0.0858 - val\_acc: 0.9798

Epoch 14/20

60000/60000 [=====] - 14s 237us/step - loss: 0.01

29 - acc: 0.9960 - val\_loss: 0.0982 - val\_acc: 0.9794

Epoch 15/20

60000/60000 [=====] - 14s 240us/step - loss: 0.01

35 - acc: 0.9960 - val\_loss: 0.1073 - val\_acc: 0.9751

Epoch 16/20

60000/60000 [=====] - 14s 238us/step - loss: 0.01

25 - acc: 0.9961 - val\_loss: 0.1004 - val\_acc: 0.9778

Epoch 17/20

60000/60000 [=====] - 14s 236us/step - loss: 0.01

40 - acc: 0.9958 - val\_loss: 0.1142 - val\_acc: 0.9771

Epoch 18/20

60000/60000 [=====] - 14s 237us/step - loss: 0.01

23 - acc: 0.9963 - val\_loss: 0.0989 - val\_acc: 0.9785

Epoch 19/20

60000/60000 [=====] - 15s 245us/step - loss: 0.00

85 - acc: 0.9979 - val\_loss: 0.1178 - val\_acc: 0.9774

Epoch 20/20

60000/60000 [=====] - 14s 237us/step - loss: 0.01

15 - acc: 0.9965 - val\_loss: 0.1039 - val\_acc: 0.9799

In [42]:

```
#Evaluat your model with accuracy and plot of (NUmber of epoches VS train_and_val_Los
s)
#Train accuracy
score = model_relu.evaluate(X_train, Y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n***** \n')
#test accuracy
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

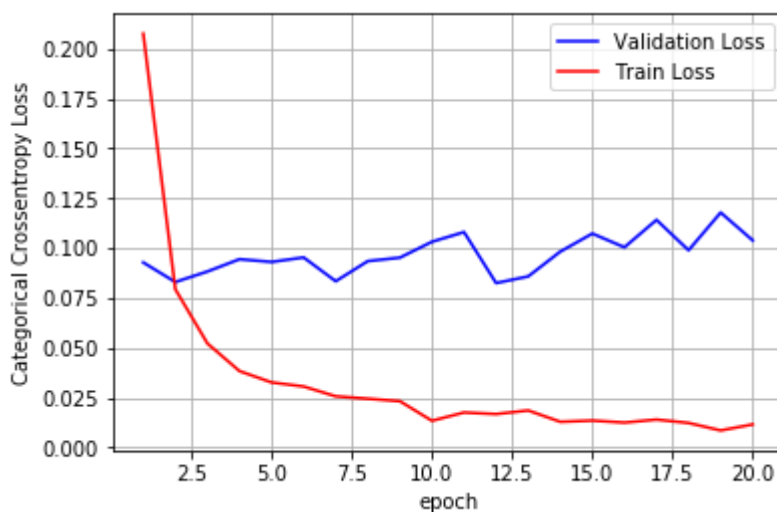
Train score: 0.011968852381023158

Train accuracy: 99.63833333333334

\*\*\*\*\*

Test score: 0.1038783551045065

Test accuracy: 97.99



In [43]:


```
w_after = model_drop.get_weights()
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)
fig = plt.figure(figsize=(15,5))
plt.title("Weight matrices after model trained")

plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 2. MLP + ReLU + adam +batch\_normalization

In [44]:

```

from keras.layers.normalization import BatchNormalization
model_batch = Sequential()
model_batch.add(Dense(610, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())
model_batch.add(Dense(420, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())
model_batch.add(Dense(210, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())
model_batch.add(Dense(output_dim, activation='softmax'))
model_batch.summary()

```

Layer (type)	Output Shape	Param #
dense_17 (Dense)	(None, 610)	478850
batch_normalization_5 (Batch Normalization)	(None, 610)	2440
dense_18 (Dense)	(None, 420)	256620
batch_normalization_6 (Batch Normalization)	(None, 420)	1680
dense_19 (Dense)	(None, 210)	88410
batch_normalization_7 (Batch Normalization)	(None, 210)	840
dense_20 (Dense)	(None, 10)	2110
Total params: 830,950		
Trainable params: 828,470		
Non-trainable params: 2,480		

In [46]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 29s 490us/step - loss: 0.19

31 - acc: 0.9414 - val\_loss: 0.1055 - val\_acc: 0.9669

Epoch 2/20

60000/60000 [=====] - 22s 373us/step - loss: 0.06

97 - acc: 0.9786 - val\_loss: 0.0878 - val\_acc: 0.9723

Epoch 3/20

60000/60000 [=====] - 20s 337us/step - loss: 0.04

55 - acc: 0.9855 - val\_loss: 0.0838 - val\_acc: 0.9722

Epoch 4/20

60000/60000 [=====] - 21s 342us/step - loss: 0.03

58 - acc: 0.9889 - val\_loss: 0.0874 - val\_acc: 0.9747

Epoch 5/20

60000/60000 [=====] - 22s 370us/step - loss: 0.02

58 - acc: 0.9919 - val\_loss: 0.0741 - val\_acc: 0.9776

Epoch 6/20

60000/60000 [=====] - 24s 394us/step - loss: 0.02

16 - acc: 0.9927 - val\_loss: 0.0791 - val\_acc: 0.9775

Epoch 7/20

60000/60000 [=====] - 21s 351us/step - loss: 0.01

94 - acc: 0.9939 - val\_loss: 0.0754 - val\_acc: 0.9791

Epoch 8/20

60000/60000 [=====] - 21s 347us/step - loss: 0.01

60 - acc: 0.9945 - val\_loss: 0.0822 - val\_acc: 0.9786

Epoch 9/20

60000/60000 [=====] - 21s 358us/step - loss: 0.01

74 - acc: 0.9941 - val\_loss: 0.0789 - val\_acc: 0.9793

Epoch 10/20

60000/60000 [=====] - 20s 334us/step - loss: 0.01

62 - acc: 0.9946 - val\_loss: 0.0899 - val\_acc: 0.9770

Epoch 11/20

60000/60000 [=====] - 20s 329us/step - loss: 0.01

09 - acc: 0.9965 - val\_loss: 0.0783 - val\_acc: 0.9797

Epoch 12/20

60000/60000 [=====] - 22s 367us/step - loss: 0.01

34 - acc: 0.9955 - val\_loss: 0.0881 - val\_acc: 0.9781

Epoch 13/20

60000/60000 [=====] - 23s 378us/step - loss: 0.01

33 - acc: 0.9955 - val\_loss: 0.0828 - val\_acc: 0.9794

Epoch 14/20

60000/60000 [=====] - 23s 390us/step - loss: 0.00

78 - acc: 0.9975 - val\_loss: 0.0667 - val\_acc: 0.9827

Epoch 15/20

60000/60000 [=====] - 21s 349us/step - loss: 0.00

80 - acc: 0.9972 - val\_loss: 0.0827 - val\_acc: 0.9800

Epoch 16/20

60000/60000 [=====] - 19s 323us/step - loss: 0.00

82 - acc: 0.9974 - val\_loss: 0.0832 - val\_acc: 0.9793

Epoch 17/20

60000/60000 [=====] - 15s 243us/step - loss: 0.00

89 - acc: 0.9971 - val\_loss: 0.0892 - val\_acc: 0.9799

Epoch 18/20

60000/60000 [=====] - 14s 228us/step - loss: 0.01

16 - acc: 0.9961 - val\_loss: 0.1042 - val\_acc: 0.9738

Epoch 19/20

60000/60000 [=====] - 15s 258us/step - loss: 0.00

86 - acc: 0.9970 - val\_loss: 0.0784 - val\_acc: 0.9809

Epoch 20/20

60000/60000 [=====] - 16s 263us/step - loss: 0.00

69 - acc: 0.9976 - val\_loss: 0.0881 - val\_acc: 0.9803



In [48]:

```
#Evaluate your model with accuracy and plot of (NUmber of epoches VS train_and_val_los
s)
#Train accuracy
score = model_batch.evaluate(X_train, Y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n***** \n')
#test accuracy
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```


Train score: 0.00419184478967436

Train accuracy: 99.87833333333333

\*\*\*\*\*

Test score: 0.08813276136659533

Test accuracy: 98.03



In [49]:

```
w_after = model_batch.get_weights()
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)
fig = plt.figure(figsize=(15,5))
plt.title("Weight matrices after model trained")

plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



### 3. MLP + ReLU + adam +dropout

In [50]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
from keras.layers import Dropout
model_drop = Sequential()
model_drop.add(Dense(610, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(420, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(210, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(output_dim, activation='softmax'))
model_drop.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_21 (Dense)	(None, 610)	478850
dropout_5 (Dropout)	(None, 610)	0
dense_22 (Dense)	(None, 420)	256620
dropout_6 (Dropout)	(None, 420)	0
dense_23 (Dense)	(None, 210)	88410
dropout_7 (Dropout)	(None, 210)	0
dense_24 (Dense)	(None, 10)	2110
=====	=====	=====
Total params: 825,990		
Trainable params: 825,990		
Non-trainable params: 0		

In [51]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 13s 222us/step - loss: 12.1  
946 - acc: 0.2356 - val\_loss: 8.1071 - val\_acc: 0.4922

Epoch 2/20

60000/60000 [=====] - 12s 197us/step - loss: 9.58  
85 - acc: 0.3991 - val\_loss: 7.1958 - val\_acc: 0.5484

Epoch 3/20

60000/60000 [=====] - 12s 200us/step - loss: 8.43  
35 - acc: 0.4718 - val\_loss: 6.4075 - val\_acc: 0.5999

Epoch 4/20

60000/60000 [=====] - 12s 196us/step - loss: 7.53  
02 - acc: 0.5274 - val\_loss: 4.6527 - val\_acc: 0.7075

Epoch 5/20

60000/60000 [=====] - 11s 185us/step - loss: 5.76  
84 - acc: 0.6372 - val\_loss: 3.8773 - val\_acc: 0.7567

Epoch 6/20

60000/60000 [=====] - 10s 175us/step - loss: 4.93  
82 - acc: 0.6893 - val\_loss: 3.2168 - val\_acc: 0.7982

Epoch 7/20

60000/60000 [=====] - 10s 174us/step - loss: 4.68  
53 - acc: 0.7055 - val\_loss: 3.1759 - val\_acc: 0.8007

Epoch 8/20

60000/60000 [=====] - 10s 174us/step - loss: 4.38  
97 - acc: 0.7244 - val\_loss: 3.0566 - val\_acc: 0.8088

Epoch 9/20

60000/60000 [=====] - 10s 174us/step - loss: 4.20  
75 - acc: 0.7361 - val\_loss: 3.2145 - val\_acc: 0.7993

Epoch 10/20

60000/60000 [=====] - 10s 175us/step - loss: 4.05  
65 - acc: 0.7455 - val\_loss: 2.8893 - val\_acc: 0.8197

Epoch 11/20

60000/60000 [=====] - 11s 184us/step - loss: 3.91  
38 - acc: 0.7547 - val\_loss: 3.0524 - val\_acc: 0.8097

Epoch 12/20

60000/60000 [=====] - 10s 174us/step - loss: 3.80  
66 - acc: 0.7616 - val\_loss: 2.8225 - val\_acc: 0.8236

Epoch 13/20

60000/60000 [=====] - 10s 174us/step - loss: 3.83  
58 - acc: 0.7599 - val\_loss: 3.0483 - val\_acc: 0.8094

Epoch 14/20

60000/60000 [=====] - 10s 174us/step - loss: 3.81  
29 - acc: 0.7613 - val\_loss: 2.8157 - val\_acc: 0.8247

Epoch 15/20

60000/60000 [=====] - 10s 174us/step - loss: 3.75  
80 - acc: 0.7652 - val\_loss: 2.8874 - val\_acc: 0.8200

Epoch 16/20

60000/60000 [=====] - 10s 174us/step - loss: 3.75  
95 - acc: 0.7652 - val\_loss: 2.8289 - val\_acc: 0.8235

Epoch 17/20

60000/60000 [=====] - 10s 174us/step - loss: 3.71  
41 - acc: 0.7681 - val\_loss: 2.9597 - val\_acc: 0.8155

Epoch 18/20

60000/60000 [=====] - 11s 177us/step - loss: 3.56  
06 - acc: 0.7777 - val\_loss: 2.8367 - val\_acc: 0.8231

Epoch 19/20

60000/60000 [=====] - 10s 174us/step - loss: 3.58  
46 - acc: 0.7763 - val\_loss: 2.8327 - val\_acc: 0.8235

Epoch 20/20

60000/60000 [=====] - 10s 173us/step - loss: 3.48  
81 - acc: 0.7821 - val\_loss: 2.7042 - val\_acc: 0.8314

In [52]:

```
#Evaluate your model with accuracy and plot of (NUmber of epoches VS train_and_val_Loss)
#Train accuracy
score = model_drop.evaluate(X_train, Y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n*****\n')
#test accuracy
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```


Train score: 2.7385955852190653

Train accuracy: 82.94666666666667

\*\*\*\*\*

Test score: 2.7042304149627685

Test accuracy: 83.14



In [53]:

```
w_after = model_drop.get_weights()
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)
fig = plt.figure(figsize=(15,5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 4. MLP + ReLU + adam +dropout+batch\_normalization

In [54]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
from keras.layers import Dropout
model_drop = Sequential()
model_drop.add(Dense(610, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(420, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(210, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(output_dim, activation='softmax'))
model_drop.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_25 (Dense)	(None, 610)	478850
batch_normalization_8 (Batch Normalization)	(None, 610)	2440
dropout_8 (Dropout)	(None, 610)	0
dense_26 (Dense)	(None, 420)	256620
batch_normalization_9 (Batch Normalization)	(None, 420)	1680
dropout_9 (Dropout)	(None, 420)	0
dense_27 (Dense)	(None, 210)	88410
batch_normalization_10 (Batch Normalization)	(None, 210)	840
dropout_10 (Dropout)	(None, 210)	0
dense_28 (Dense)	(None, 10)	2110
=====	=====	=====
Total params: 830,950		
Trainable params: 828,470		
Non-trainable params: 2,480		



In [55]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 17s 279us/step - loss: 0.67  
29 - acc: 0.7925 - val\_loss: 0.1967 - val\_acc: 0.9368

Epoch 2/20

60000/60000 [=====] - 14s 236us/step - loss: 0.30  
22 - acc: 0.9089 - val\_loss: 0.1463 - val\_acc: 0.9517

Epoch 3/20

60000/60000 [=====] - 13s 219us/step - loss: 0.23  
94 - acc: 0.9295 - val\_loss: 0.1262 - val\_acc: 0.9610

Epoch 4/20

60000/60000 [=====] - 13s 225us/step - loss: 0.20  
20 - acc: 0.9397 - val\_loss: 0.1061 - val\_acc: 0.9669

Epoch 5/20

60000/60000 [=====] - 13s 212us/step - loss: 0.17  
87 - acc: 0.9465 - val\_loss: 0.0994 - val\_acc: 0.9705

Epoch 6/20

60000/60000 [=====] - 14s 234us/step - loss: 0.16  
23 - acc: 0.9520 - val\_loss: 0.0983 - val\_acc: 0.9719

Epoch 7/20

60000/60000 [=====] - 14s 226us/step - loss: 0.14  
73 - acc: 0.9554 - val\_loss: 0.0881 - val\_acc: 0.9736

Epoch 8/20

60000/60000 [=====] - 15s 251us/step - loss: 0.13  
89 - acc: 0.9581 - val\_loss: 0.0823 - val\_acc: 0.9742

Epoch 9/20

60000/60000 [=====] - 19s 312us/step - loss: 0.13  
06 - acc: 0.9610 - val\_loss: 0.0802 - val\_acc: 0.9756

Epoch 10/20

60000/60000 [=====] - 17s 288us/step - loss: 0.12  
18 - acc: 0.9628 - val\_loss: 0.0745 - val\_acc: 0.9762

Epoch 11/20

60000/60000 [=====] - 19s 320us/step - loss: 0.11  
33 - acc: 0.9665 - val\_loss: 0.0711 - val\_acc: 0.9792

Epoch 12/20

60000/60000 [=====] - 19s 314us/step - loss: 0.10  
91 - acc: 0.9669 - val\_loss: 0.0732 - val\_acc: 0.9781

Epoch 13/20

60000/60000 [=====] - 18s 295us/step - loss: 0.10  
52 - acc: 0.9679 - val\_loss: 0.0660 - val\_acc: 0.9807

Epoch 14/20

60000/60000 [=====] - 17s 282us/step - loss: 0.10  
00 - acc: 0.9695 - val\_loss: 0.0660 - val\_acc: 0.9793

Epoch 15/20

60000/60000 [=====] - 16s 266us/step - loss: 0.09  
54 - acc: 0.9712 - val\_loss: 0.0694 - val\_acc: 0.9802

Epoch 16/20

60000/60000 [=====] - 16s 272us/step - loss: 0.09  
01 - acc: 0.9729 - val\_loss: 0.0635 - val\_acc: 0.9825

Epoch 17/20

60000/60000 [=====] - 17s 280us/step - loss: 0.08  
25 - acc: 0.9746 - val\_loss: 0.0609 - val\_acc: 0.9832

Epoch 18/20

60000/60000 [=====] - 19s 310us/step - loss: 0.08  
02 - acc: 0.9756 - val\_loss: 0.0636 - val\_acc: 0.9819

Epoch 19/20

60000/60000 [=====] - 20s 336us/step - loss: 0.08  
11 - acc: 0.9752 - val\_loss: 0.0622 - val\_acc: 0.9820

Epoch 20/20

60000/60000 [=====] - 17s 283us/step - loss: 0.07  
69 - acc: 0.9763 - val\_loss: 0.0628 - val\_acc: 0.9821

In [56]:

```
#Evaluate your model with accuracy and plot of (NUmber of epoches VS train_and_val_los
s)
#Train accuracy
score = model_drop.evaluate(X_train, Y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n*****\n')
#test accuracy
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```


Train score: 0.02124301015394934

Train accuracy: 99.33166666666666

\*\*\*\*\*

Test score: 0.0627561581715534

Test accuracy: 98.21




In [57]:

```
w_after = model_drop.get_weights()
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)
fig = plt.figure(figsize=(15,5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## Model 3 -- with 5 Hidden layers

### 1. MLP + ReLU + adam

In [58]:

```

# Multilayer perceptron
# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 
model_relu = Sequential()
model_relu.add(Dense(690, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(530, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(412, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(231, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))
model_relu.summary()

```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_29 (Dense)	(None, 690)	541650
dense_30 (Dense)	(None, 530)	366230
dense_31 (Dense)	(None, 412)	218772
dense_32 (Dense)	(None, 231)	95403
dense_33 (Dense)	(None, 112)	25984
dense_34 (Dense)	(None, 10)	1130
=====	=====	=====
Total params: 1,249,169		
Trainable params: 1,249,169		
Non-trainable params: 0		

In [59]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 17s 290us/step - loss: 0.28  
13 - acc: 0.9226 - val\_loss: 0.1477 - val\_acc: 0.9539

Epoch 2/20

60000/60000 [=====] - 15s 249us/step - loss: 0.10  
31 - acc: 0.9676 - val\_loss: 0.1105 - val\_acc: 0.9676

Epoch 3/20

60000/60000 [=====] - 17s 281us/step - loss: 0.07  
02 - acc: 0.9777 - val\_loss: 0.1078 - val\_acc: 0.9687

Epoch 4/20

60000/60000 [=====] - 17s 283us/step - loss: 0.05  
54 - acc: 0.9823 - val\_loss: 0.0959 - val\_acc: 0.9752

Epoch 5/20

60000/60000 [=====] - 16s 259us/step - loss: 0.04  
81 - acc: 0.9847 - val\_loss: 0.0991 - val\_acc: 0.9724

Epoch 6/20

60000/60000 [=====] - 17s 277us/step - loss: 0.04  
09 - acc: 0.9870 - val\_loss: 0.0890 - val\_acc: 0.9738

Epoch 7/20

60000/60000 [=====] - 15s 258us/step - loss: 0.03  
54 - acc: 0.9890 - val\_loss: 0.1293 - val\_acc: 0.9684

Epoch 8/20

60000/60000 [=====] - 15s 256us/step - loss: 0.03  
70 - acc: 0.9884 - val\_loss: 0.0868 - val\_acc: 0.9764

Epoch 9/20

60000/60000 [=====] - 17s 289us/step - loss: 0.03  
28 - acc: 0.9897 - val\_loss: 0.0954 - val\_acc: 0.9750

Epoch 10/20

60000/60000 [=====] - 16s 265us/step - loss: 0.02  
63 - acc: 0.9915 - val\_loss: 0.0797 - val\_acc: 0.9799

Epoch 11/20

60000/60000 [=====] - 17s 289us/step - loss: 0.02  
52 - acc: 0.9925 - val\_loss: 0.1105 - val\_acc: 0.9750

Epoch 12/20

60000/60000 [=====] - 16s 267us/step - loss: 0.02  
41 - acc: 0.9925 - val\_loss: 0.0822 - val\_acc: 0.9801

Epoch 13/20

60000/60000 [=====] - 18s 295us/step - loss: 0.02  
07 - acc: 0.9937 - val\_loss: 0.0962 - val\_acc: 0.9777

Epoch 14/20

60000/60000 [=====] - 18s 306us/step - loss: 0.02  
29 - acc: 0.9931 - val\_loss: 0.0887 - val\_acc: 0.9770

Epoch 15/20

60000/60000 [=====] - 19s 322us/step - loss: 0.01  
77 - acc: 0.9942 - val\_loss: 0.0838 - val\_acc: 0.9815

Epoch 16/20

60000/60000 [=====] - 17s 275us/step - loss: 0.01  
66 - acc: 0.9952 - val\_loss: 0.1196 - val\_acc: 0.9735

Epoch 17/20

60000/60000 [=====] - 17s 278us/step - loss: 0.01  
86 - acc: 0.9946 - val\_loss: 0.1064 - val\_acc: 0.9767

Epoch 18/20

60000/60000 [=====] - 18s 295us/step - loss: 0.01  
82 - acc: 0.9946 - val\_loss: 0.0977 - val\_acc: 0.9782

Epoch 19/20

60000/60000 [=====] - 19s 317us/step - loss: 0.01  
41 - acc: 0.9960 - val\_loss: 0.1035 - val\_acc: 0.9790

Epoch 20/20

60000/60000 [=====] - 18s 296us/step - loss: 0.01  
38 - acc: 0.9963 - val\_loss: 0.0970 - val\_acc: 0.9802

In [60]:

```
#Evaluate your model with accuracy and plot of (NUmber of epoches VS train_and_val_Los
s)
#Train accuracy
score = model_relu.evaluate(X_train, Y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n*****\n')
#test accuracy
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```


Train score: 0.007615429904182383

Train accuracy: 99.78333333333333

\*\*\*\*\*

Test score: 0.09702783975718078

Test accuracy: 98.02





In [61]:

```
w_after = model_relu.get_weights()
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)
fig = plt.figure(figsize=(15,5))
plt.title("Weight matrices after model trained")

plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w,color='g')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w,color='b')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 2. MLP + ReLU + adam +batch\_normalization

In [62]:

```
# Multilayer perceptron
# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(\theta, \sigma) = N(\theta, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(\theta, \sigma) = N(\theta, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 
model_relu = Sequential()
model_relu.add(Dense(690, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(530, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(BatchNormalization())
model_relu.add(Dense(412, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(BatchNormalization())
model_relu.add(Dense(231, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(BatchNormalization())
model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))
model_relu.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_35 (Dense)	(None, 690)	541650
batch_normalization_11 (Batch Normalization)	(None, 690)	2760
dense_36 (Dense)	(None, 530)	366230
batch_normalization_12 (Batch Normalization)	(None, 530)	2120
dense_37 (Dense)	(None, 412)	218772
batch_normalization_13 (Batch Normalization)	(None, 412)	1648
dense_38 (Dense)	(None, 231)	95403
batch_normalization_14 (Batch Normalization)	(None, 231)	924
dense_39 (Dense)	(None, 112)	25984
batch_normalization_15 (Batch Normalization)	(None, 112)	448
dense_40 (Dense)	(None, 10)	1130
=====	=====	=====
Total params: 1,257,069		
Trainable params: 1,253,119		
Non-trainable params: 3,950		

In [63]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 23s 383us/step - loss: 0.20

23 - acc: 0.9385 - val\_loss: 0.1073 - val\_acc: 0.9675

Epoch 2/20

60000/60000 [=====] - 18s 305us/step - loss: 0.07

37 - acc: 0.9770 - val\_loss: 0.0832 - val\_acc: 0.9740

Epoch 3/20

60000/60000 [=====] - 20s 327us/step - loss: 0.05

62 - acc: 0.9816 - val\_loss: 0.0765 - val\_acc: 0.9750

Epoch 4/20

60000/60000 [=====] - 21s 344us/step - loss: 0.03

84 - acc: 0.9880 - val\_loss: 0.0715 - val\_acc: 0.9772

Epoch 5/20

60000/60000 [=====] - 20s 333us/step - loss: 0.03

67 - acc: 0.9874 - val\_loss: 0.0700 - val\_acc: 0.9789

Epoch 6/20

60000/60000 [=====] - 18s 297us/step - loss: 0.03

05 - acc: 0.9901 - val\_loss: 0.0841 - val\_acc: 0.9759

Epoch 7/20

60000/60000 [=====] - 19s 313us/step - loss: 0.03

15 - acc: 0.9895 - val\_loss: 0.0850 - val\_acc: 0.9780

Epoch 8/20

60000/60000 [=====] - 19s 318us/step - loss: 0.02

58 - acc: 0.9914 - val\_loss: 0.0867 - val\_acc: 0.9753

Epoch 9/20

60000/60000 [=====] - 19s 321us/step - loss: 0.02

46 - acc: 0.9919 - val\_loss: 0.0757 - val\_acc: 0.9785

Epoch 10/20

60000/60000 [=====] - 20s 327us/step - loss: 0.02

01 - acc: 0.9935 - val\_loss: 0.0708 - val\_acc: 0.9793

Epoch 11/20

60000/60000 [=====] - 20s 337us/step - loss: 0.02

14 - acc: 0.9926 - val\_loss: 0.0746 - val\_acc: 0.9795

Epoch 12/20

60000/60000 [=====] - 21s 350us/step - loss: 0.01

98 - acc: 0.9934 - val\_loss: 0.0649 - val\_acc: 0.9814

Epoch 13/20

60000/60000 [=====] - 22s 367us/step - loss: 0.01

79 - acc: 0.9941 - val\_loss: 0.0792 - val\_acc: 0.9779

Epoch 14/20

60000/60000 [=====] - 21s 347us/step - loss: 0.01

71 - acc: 0.9943 - val\_loss: 0.0770 - val\_acc: 0.9804

Epoch 15/20

60000/60000 [=====] - 19s 323us/step - loss: 0.01

50 - acc: 0.9950 - val\_loss: 0.0813 - val\_acc: 0.9801

Epoch 16/20

60000/60000 [=====] - 16s 268us/step - loss: 0.01

17 - acc: 0.9965 - val\_loss: 0.0822 - val\_acc: 0.9792

Epoch 17/20

60000/60000 [=====] - 16s 267us/step - loss: 0.01

57 - acc: 0.9947 - val\_loss: 0.0753 - val\_acc: 0.9804

Epoch 18/20

60000/60000 [=====] - 16s 265us/step - loss: 0.01

29 - acc: 0.9956 - val\_loss: 0.0876 - val\_acc: 0.9790

Epoch 19/20

60000/60000 [=====] - 16s 263us/step - loss: 0.01

20 - acc: 0.9960 - val\_loss: 0.0897 - val\_acc: 0.9780

Epoch 20/20

60000/60000 [=====] - 16s 263us/step - loss: 0.01

27 - acc: 0.9959 - val\_loss: 0.0660 - val\_acc: 0.9834

In [64]:

```

#Evaluate your model with accuracy and plot of (NUmber of epochs VS train_and_val_los
s)
#Train accuracy
score = model_relu.evaluate(X_train, Y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n*****\n')
#test accuracy
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```


Train score: 0.006093753856421487

Train accuracy: 99.79833333333333

\*\*\*\*\*

Test score: 0.06601150656397804

Test accuracy: 98.34



In [65]:

```
w_after = model_relu.get_weights()
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)
fig = plt.figure(figsize=(15,5))
plt.title("Weight matrices after model trained")

plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w,color='g')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w,color='b')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



### 3. MLP + ReLU + adam + dropout

In [66]:

```
# Multilayer perceptron
# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 
model_relu = Sequential()
model_relu.add(Dense(690, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
#model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(530, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
#model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(412, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
#model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(231, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
#model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
#model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(output_dim, activation='softmax'))
model_relu.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_41 (Dense)	(None, 690)	541650
dropout_11 (Dropout)	(None, 690)	0
dense_42 (Dense)	(None, 530)	366230
dropout_12 (Dropout)	(None, 530)	0
dense_43 (Dense)	(None, 412)	218772
dropout_13 (Dropout)	(None, 412)	0
dense_44 (Dense)	(None, 231)	95403
dropout_14 (Dropout)	(None, 231)	0
dense_45 (Dense)	(None, 112)	25984
dropout_15 (Dropout)	(None, 112)	0
dense_46 (Dense)	(None, 10)	1130
=====	=====	=====
Total params: 1,249,169		
Trainable params: 1,249,169		
Non-trainable params: 0		



In [67]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 17s 282us/step - loss: 8.82  
55 - acc: 0.2601 - val\_loss: 1.2662 - val\_acc: 0.6625

Epoch 2/20

60000/60000 [=====] - 14s 240us/step - loss: 1.32  
11 - acc: 0.5528 - val\_loss: 0.7544 - val\_acc: 0.7617

Epoch 3/20

60000/60000 [=====] - 14s 241us/step - loss: 0.86  
94 - acc: 0.7162 - val\_loss: 0.5276 - val\_acc: 0.8670

Epoch 4/20

60000/60000 [=====] - 14s 240us/step - loss: 0.61  
93 - acc: 0.8110 - val\_loss: 0.3173 - val\_acc: 0.9238

Epoch 5/20

60000/60000 [=====] - 15s 243us/step - loss: 0.48  
21 - acc: 0.8656 - val\_loss: 0.2496 - val\_acc: 0.9372

Epoch 6/20

60000/60000 [=====] - 14s 241us/step - loss: 0.40  
34 - acc: 0.8948 - val\_loss: 0.2200 - val\_acc: 0.9475

Epoch 7/20

60000/60000 [=====] - 14s 240us/step - loss: 0.34  
84 - acc: 0.9117 - val\_loss: 0.1971 - val\_acc: 0.9506

Epoch 8/20

60000/60000 [=====] - 14s 240us/step - loss: 0.31  
98 - acc: 0.9193 - val\_loss: 0.1902 - val\_acc: 0.9541

Epoch 9/20

60000/60000 [=====] - 15s 243us/step - loss: 0.28  
56 - acc: 0.9282 - val\_loss: 0.1759 - val\_acc: 0.9565

Epoch 10/20

60000/60000 [=====] - 14s 241us/step - loss: 0.26  
54 - acc: 0.9344 - val\_loss: 0.1572 - val\_acc: 0.9599

Epoch 11/20

60000/60000 [=====] - 14s 240us/step - loss: 0.24  
93 - acc: 0.9382 - val\_loss: 0.1440 - val\_acc: 0.9655

Epoch 12/20

60000/60000 [=====] - 14s 241us/step - loss: 0.23  
19 - acc: 0.9433 - val\_loss: 0.1354 - val\_acc: 0.9663

Epoch 13/20

60000/60000 [=====] - 14s 241us/step - loss: 0.22  
19 - acc: 0.9452 - val\_loss: 0.1430 - val\_acc: 0.9661

Epoch 14/20

60000/60000 [=====] - 14s 242us/step - loss: 0.21  
14 - acc: 0.9490 - val\_loss: 0.1399 - val\_acc: 0.9655

Epoch 15/20

60000/60000 [=====] - 14s 241us/step - loss: 0.19  
60 - acc: 0.9513 - val\_loss: 0.1356 - val\_acc: 0.9666

Epoch 16/20

60000/60000 [=====] - 14s 241us/step - loss: 0.18  
86 - acc: 0.9537 - val\_loss: 0.1317 - val\_acc: 0.9679

Epoch 17/20

60000/60000 [=====] - 14s 241us/step - loss: 0.18  
67 - acc: 0.9545 - val\_loss: 0.1302 - val\_acc: 0.9678

Epoch 18/20

60000/60000 [=====] - 14s 241us/step - loss: 0.17  
29 - acc: 0.9578 - val\_loss: 0.1254 - val\_acc: 0.9690

Epoch 19/20

60000/60000 [=====] - 14s 240us/step - loss: 0.16  
58 - acc: 0.9609 - val\_loss: 0.1231 - val\_acc: 0.9711

Epoch 20/20

60000/60000 [=====] - 15s 244us/step - loss: 0.15  
90 - acc: 0.9606 - val\_loss: 0.1188 - val\_acc: 0.9736

In [68]:

```
#Evaluat your model with accuracy and plot of (NUmber of epochs VS train_and_val_Los
s)
#Train accuracy
score = model_relu.evaluate(X_train, Y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n***** \n')
#test accuracy
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# List of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```


Train score: 0.06144777707796311

Train accuracy: 98.44166666666668

\*\*\*\*\*

Test score: 0.11881388411391526

Test accuracy: 97.36



In [70]:

```
w_after = model_relu.get_weights()
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)
fig = plt.figure(figsize=(15,5))
plt.title("Weight matrices after model trained")

plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w,color='g')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w,color='b')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 4. MLP + ReLU + adam + dropout + batch\_normalization

In [71]:

```
# Multilayer perceptron
# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 
model_relu = Sequential()
model_relu.add(Dense(690, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(530, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(412, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(231, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(output_dim, activation='softmax'))
model_relu.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_47 (Dense)	(None, 690)	541650
batch_normalization_16 (Batch Normalization)	(None, 690)	2760
dropout_16 (Dropout)	(None, 690)	0
dense_48 (Dense)	(None, 530)	366230
batch_normalization_17 (Batch Normalization)	(None, 530)	2120
dropout_17 (Dropout)	(None, 530)	0
dense_49 (Dense)	(None, 412)	218772
batch_normalization_18 (Batch Normalization)	(None, 412)	1648
dropout_18 (Dropout)	(None, 412)	0
dense_50 (Dense)	(None, 231)	95403
batch_normalization_19 (Batch Normalization)	(None, 231)	924
dropout_19 (Dropout)	(None, 231)	0
dense_51 (Dense)	(None, 112)	25984
batch_normalization_20 (Batch Normalization)	(None, 112)	448
dropout_20 (Dropout)	(None, 112)	0
dense_52 (Dense)	(None, 10)	1130
=====	=====	=====
Total params: 1,257,069		
Trainable params: 1,253,119		
Non-trainable params: 3,950		
=====		

In [72]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```



Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 26s 428us/step - loss: 1.05  
27 - acc: 0.6697 - val\_loss: 0.2500 - val\_acc: 0.9261

Epoch 2/20

60000/60000 [=====] - 21s 357us/step - loss: 0.36  
79 - acc: 0.8910 - val\_loss: 0.1626 - val\_acc: 0.9519

Epoch 3/20

60000/60000 [=====] - 24s 406us/step - loss: 0.26  
47 - acc: 0.9257 - val\_loss: 0.1312 - val\_acc: 0.9607

Epoch 4/20

60000/60000 [=====] - 24s 404us/step - loss: 0.22  
13 - acc: 0.9374 - val\_loss: 0.1179 - val\_acc: 0.9667

Epoch 5/20

60000/60000 [=====] - 23s 385us/step - loss: 0.19  
19 - acc: 0.9465 - val\_loss: 0.1039 - val\_acc: 0.9714

Epoch 6/20

60000/60000 [=====] - 24s 407us/step - loss: 0.17  
00 - acc: 0.9528 - val\_loss: 0.0911 - val\_acc: 0.9737

Epoch 7/20

60000/60000 [=====] - 22s 375us/step - loss: 0.15  
57 - acc: 0.9561 - val\_loss: 0.0951 - val\_acc: 0.9737

Epoch 8/20

60000/60000 [=====] - 21s 346us/step - loss: 0.14  
61 - acc: 0.9586 - val\_loss: 0.0824 - val\_acc: 0.9764

Epoch 9/20

60000/60000 [=====] - 21s 353us/step - loss: 0.12  
85 - acc: 0.9640 - val\_loss: 0.0786 - val\_acc: 0.9789

Epoch 10/20

60000/60000 [=====] - 21s 356us/step - loss: 0.12  
62 - acc: 0.9649 - val\_loss: 0.0752 - val\_acc: 0.9795

Epoch 11/20

60000/60000 [=====] - 23s 391us/step - loss: 0.11  
90 - acc: 0.9670 - val\_loss: 0.0717 - val\_acc: 0.9807

Epoch 12/20

60000/60000 [=====] - 25s 414us/step - loss: 0.11  
07 - acc: 0.9692 - val\_loss: 0.0718 - val\_acc: 0.9806

Epoch 13/20

60000/60000 [=====] - 20s 333us/step - loss: 0.10  
66 - acc: 0.9704 - val\_loss: 0.0714 - val\_acc: 0.9812

Epoch 14/20

60000/60000 [=====] - 22s 365us/step - loss: 0.10  
21 - acc: 0.9717 - val\_loss: 0.0789 - val\_acc: 0.9791

Epoch 15/20

60000/60000 [=====] - 20s 341us/step - loss: 0.09  
61 - acc: 0.9729 - val\_loss: 0.0664 - val\_acc: 0.9830

Epoch 16/20

60000/60000 [=====] - 25s 412us/step - loss: 0.09  
52 - acc: 0.9736 - val\_loss: 0.0634 - val\_acc: 0.9838

Epoch 17/20

60000/60000 [=====] - 24s 398us/step - loss: 0.08  
79 - acc: 0.9755 - val\_loss: 0.0694 - val\_acc: 0.9835

Epoch 18/20

60000/60000 [=====] - 20s 335us/step - loss: 0.08  
39 - acc: 0.9762 - val\_loss: 0.0666 - val\_acc: 0.9824

Epoch 19/20

60000/60000 [=====] - 20s 337us/step - loss: 0.08  
53 - acc: 0.9760 - val\_loss: 0.0633 - val\_acc: 0.9829

Epoch 20/20

60000/60000 [=====] - 20s 341us/step - loss: 0.08  
04 - acc: 0.9769 - val\_loss: 0.0630 - val\_acc: 0.9830

In [73]:

```
#Evaluate your model with accuracy and plot of (NUmber of epochs VS train_and_val_loss)
#Train accuracy
score = model_relu.evaluate(X_train, Y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n***** \n')
#test accuracy
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# List of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```


Train score: 0.019241652972002823

Train accuracy: 99.47833333333334

\*\*\*\*\*

Test score: 0.06302054594261572

Test accuracy: 98.3



In [74]:

```
w_after = model_relu.get_weights()
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)
fig = plt.figure(figsize=(15,5))
plt.title("Weight matrices after model trained")

plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w,color='g')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w,color='b')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## CONCLUSION:

In [76]:

```
from prettytable import PrettyTable
tb = PrettyTable()
tb.field_names= ("Hidden Layers", "Model", "Accuracy")
tb.add_row(["2", "MLP + ADAM + RELU",97.82])
tb.add_row(["2", "MLP + ADAM + RELU + batch_normalization",98.2])
tb.add_row(["2", "MLP + ADAM + RELU + dropout",97.77])
tb.add_row(["2", "MLP + ADAM + RELU + dropout+ batch_normalization",98.19])
tb.add_row([" ", " ", " "])
tb.add_row([" ", " ", " "])
tb.field_names= ("Hidden Layers", "Model", "Accuracy")
tb.add_row(["3", "MLP + ADAM + RELU",97.99])
tb.add_row(["3", "MLP + ADAM + RELU + batch_normalization",98.03])
tb.add_row(["3", "MLP + ADAM + RELU + dropout",83.14])
tb.add_row(["3", "MLP + ADAM + RELU + dropout+ batch_normalization",98.21])
tb.add_row([" ", " ", " "])
tb.add_row([" ", " ", " "])
tb.field_names= ("Hidden Layers", "Model", "Accuracy")
tb.add_row(["5", "MLP + ADAM + RELU",98.02])
tb.add_row(["5", "MLP + ADAM + RELU + batch_normalization",98.34])
tb.add_row(["5", "MLP + ADAM + RELU + dropout",97.36])
tb.add_row(["5", "MLP + ADAM + RELU + dropout+ batch_normalization",98.3])
print(tb.get_string(titles = "MLP Models - Observations"))
```

+-----+-----+-----+		
-----+		
Hidden Layers	Model	Accur
acy		
+-----+-----+-----+		
-----+		
2	MLP + ADAM + RELU	97.8
2		
2	MLP + ADAM + RELU + batch_normalization	98.
2		
2	MLP + ADAM + RELU + dropout	97.7
7		
2	MLP + ADAM + RELU + dropout+ batch_normalization	98.1
9		
3	MLP + ADAM + RELU	97.9
9		
3	MLP + ADAM + RELU + batch_normalization	98.0
3		
3	MLP + ADAM + RELU + dropout	83.1
4		
3	MLP + ADAM + RELU + dropout+ batch_normalization	98.2
1		
5	MLP + ADAM + RELU	98.0
2		
5	MLP + ADAM + RELU + batch_normalization	98.3
4		
5	MLP + ADAM + RELU + dropout	97.3
6		
5	MLP + ADAM + RELU + dropout+ batch_normalization	98.
3		
+-----+-----+-----+		
-----+		